

Quick Intro to Debugging with GDB


Christian A. Pagot



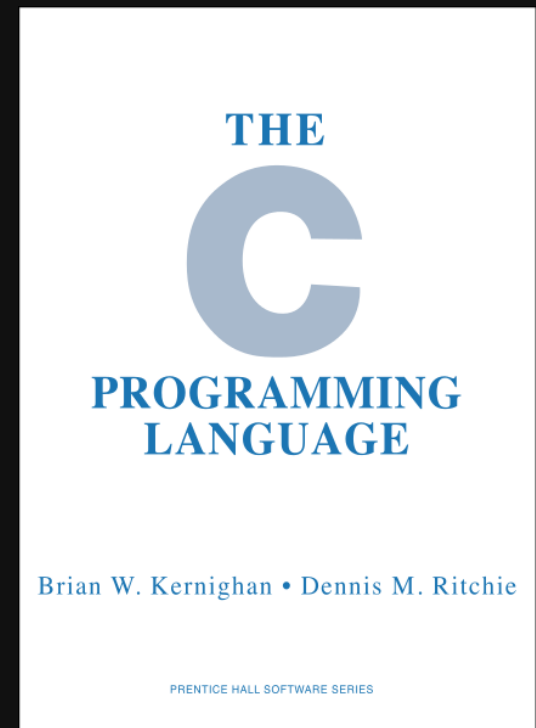
Universidade Federal da Paraíba
Centro de Informática

What is C?

C is a **imperative, general-purpose** programming language, developed in 1973 at AT&T Bell Labs by **Ken Thompson** (left) and **Dennis Ritchie** (right):



unknown



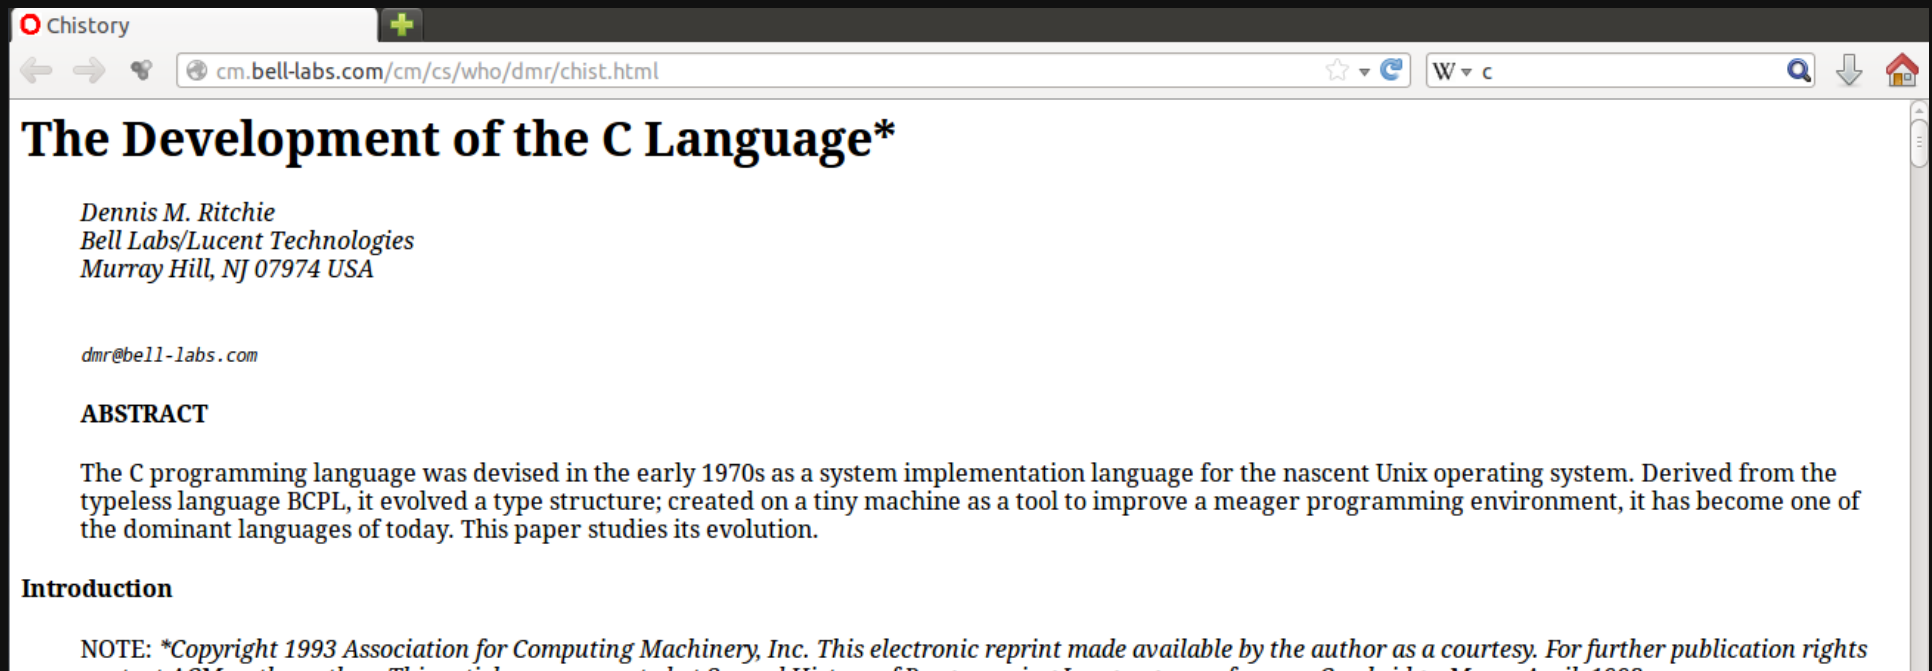
Julesmazur (Wikipedia)



The Development of the C Language

“*The Development of the C Language*”, by Dennis Ritchie:

- <http://cm.bell-labs.com/cm/cs/who/dmr/chist.html>



The screenshot shows a web browser window with the title bar 'Chistory'. The address bar contains the URL 'cm.bell-labs.com/cm/cs/who/dmr/chist.html'. The page content is titled 'The Development of the C Language*' and is attributed to 'Dennis M. Ritchie, Bell Labs/Lucent Technologies, Murray Hill, NJ 07974 USA'. The email address 'dmr@bell-labs.com' is listed. The 'ABSTRACT' section states: 'The C programming language was devised in the early 1970s as a system implementation language for the nascent Unix operating system. Derived from the typeless language BCPL, it evolved a type structure; created on a tiny machine as a tool to improve a meager programming environment, it has become one of the dominant languages of today. This paper studies its evolution.' The 'Introduction' section begins with a note: '*Copyright 1993 Association for Computing Machinery, Inc. This electronic reprint made available by the author as a courtesy. For further publication rights contact ACM.' The bottom of the page contains a footer with the logo of the Universidade Federal da Paraíba and the text 'Centro de Informática'.

Chistory

cm.bell-labs.com/cm/cs/who/dmr/chist.html

The Development of the C Language*

Dennis M. Ritchie
Bell Labs/Lucent Technologies
Murray Hill, NJ 07974 USA

dmr@bell-labs.com

ABSTRACT

The C programming language was devised in the early 1970s as a system implementation language for the nascent Unix operating system. Derived from the typeless language BCPL, it evolved a type structure; created on a tiny machine as a tool to improve a meager programming environment, it has become one of the dominant languages of today. This paper studies its evolution.

Introduction

NOTE: *Copyright 1993 Association for Computing Machinery, Inc. This electronic reprint made available by the author as a courtesy. For further publication rights contact ACM.



Features of the C Language

- C is **imperative** (procedural).

Instructions define the actions of the processor.

- C instructions **map easily** to **machine instructions**.

- C is meant to be **cross-platform**.

Source code can be compiled to different hardware with minimal modifications.

- C uses **lexical scoping**.

Variable scope defined by its position in the source code.



Features of the C Language

- C uses **static** type system.

Types are checked during compile time.

- C supports **recursion**.

A function can call itself.

- In C, function parameters are **passed by value**.

Copies.

- C is **weakly** typed.

Casting.



C Compilers

- There are several **C compilers** around. Some examples are:
 - Open Watcom C/C++.
 - CodeWarrior.
 - Clang (LLVM).
 - **GCC**.
- More C compilers can be found on:
https://en.wikipedia.org/wiki/List_of_compilers#C_compilers



Versions of C

- Since its creation, C has undergone several **improvements**.
- Resulting **versions** have been **published** as **standards**:

- ANSI C (1990), ratified as ISO/IEC 9899:1990.

`gcc` option: `-ansi` or `-std=c90`

- ISO/IEC 9899:1999.

`gcc` option: `-std=c99`

- ISO/IEC 9899:2011.

`gcc` option: `-std=c11`

At the time of this writing, **my gcc defaults** to `-std=gnu89`, which stands for the **GNU-extended** version of the **ISO/IEC 9899:1990**.



C Program Example

The program below computes the summation of all integers within a given closed interval:

summation.c

```
#include <stdio.h>

int Sum( int begin, int end ) {
    int i;
    int acc = 0;
    for ( i = begin; i <= end; i++ )
        acc += i;

    return acc;
}

int main ( void ) {
    int a = 1;
    int b = 5;
    int sum = Sum( a, b );
    printf( "Sum: %i\n", sum );
    return 0;
}
```

Compile and run

```
~$ gcc summation.c
~$ ./a.out
```

or

```
~$ gcc summation.c -o summation
~$ ./summation
```

DIY!



C Data Types

- Basic types.
- Structured types.
- User defined types.

They are, basically, combinations of the above data types identified by a user-defined name.



Basic Data Types

- char
- int
- float
- double
- pointers

Optional specifiers*:

- signed
- unsigned
- short
- long

*may not apply to all numeric types.

```
char
signed char
unsigned char
short
short int
signed short
signed short int
unsigned short
unsigned short int
int
```

```
signed int
unsigned
unsigned int
long
long int
signed long
signed long int
unsigned long
unsigned long int
long long
```

```
long long int
signed long long
signed long long int
unsigned long long
unsigned long long int
float
double
long double
```



Aggregate Data Types

C offers 2 types of aggregate data types:

- Arrays.
- Structs.



Structs

- Is an **aggregate** that might contain **members** of **different types**.
- Given a variable of type `struct`, its members can be **accessed** through the `'.'` operator.

- Example

`structs.c`

```
struct Date {  
    int day;  
    int month;  
    int year;  
};  
  
int main( void ) {  
    struct Date x;  
    x.day = 5;  
    x.month = 2;  
    x.year = 2016;  
  
    return 0;  
}
```



Arrays

*“(...) an **array data structure**, or simply an **array**, is a data structure consisting of a **collection** of elements (values or variables), each identified by at least one array **index** or key. An array is stored so that the **position** of each element **can be computed from its index tuple** by a mathematical formula.”*

Array data structure, Wikipedia



Arrays in C

- Traditionally of **fixed, static** size.
- Usually, **all elements** are of the **same type**.
- **Does not** carry information about its **size**!
- May be **multidimensional**.
- Example:

example_26.c

```
#include <stdio.h>

int x[4] = { 10, 20, 30, 40 };

int main( void ) {
    int i;
    for ( i = 0; i < 4; i++ )
        printf( "%i ", x[i] );

    return 0;
}
```



Multidimensional Arrays in C

- **Multidimensional** arrays can be obtained by applying the array concept **recursively**.
- **Example:**

example_27.c

```
#include <stdio.h>

int x[2][3] = { { 10, 20, 30}, { 40, 50, 60 } };

int main( void ) {
    int i, j;
    for ( i = 0; i < 2; i++ )
        for ( j = 0; j < 3; j++ )
            printf( "%i ", x[i][j] );

    return 0;
}
```

Logical representation
of the 2D array.



	0	1	2
0	10	20	30
1	40	50	60

addr + 0	10
addr + 1	20
addr + 2	30
addr + 3	40
addr + 4	50
addr + 5	60

**Actual distribution of the
array elements in memory.**



Another C Program Example

The program below computes the average of the following integers: 2, 2, 3, 4 and 5:

average.c

```
#include <stdio.h>

float Average( int *w, int n ) {
    int i;
    int avg;
    for ( i = 0; i < n; i++ )
        avg += w[i];

    return avg / n;
}

int main ( void ) {
    int x[5] = {2, 2, 3, 4, 5};
    int num = 5;
    float avg = Average( x, num );
    printf( "Average: %f\n", avg );
    return 0;
}
```

Oops! We were expecting the average to be 3.2! How do we approach this problem?

DIY!



Debuggers

- “(...) a computer program that is used to test and debug other programs (...).”

Debugger, Wikipedia.


- Examples:
 - Microsoft Visual Studio Debugger.
 - LLDB.
 - GDB.



GDB

GDB, or the GNU Project Debugger

- Popular debugger tool used among Unix/Linux programmers.
- It comes, usually, pre-installed in several Linux distributions.



The Archer Fish,
the GDB mascot.



GDB

Features (as of version 7.10):

- Can be used to debug C and C++ programs.
- Partial support to some other languages.
- Text-based.
- “Normal”, temporary and conditional breakpoints.
- Single-stepping.
- Resume.
- Watchpoints.
- Variable inspection.
- Call stack inspection.
- Etc.



GDB Usage Example

Back to our broken C program:

average.c

```
#include <stdio.h>

float Average( int *w, int n ) {
    int i;
    int avg;
    for ( i = 0; i < n; i++ )
        avg += w[i];

    return avg / n;
}

int main ( void ) {
    int x[5] = {2, 2, 3, 4, 5};
    int num = 5;
    float avg = Average( x, num );
    printf( "Average: %f\n", avg );
    return 0;
}
```

Recompile the program with the following command:

```
~$ gcc -g3 average.c -o average
```

Inserts debugging information into the executable.



GDB Usage Example

- After the recompilation, invoke GDB on the executable file:

```
~$ gdb ./average
```

- From within the GDB environment, let's issue the following commands:

- List the source code from within GDB:

```
(gdb) list 1, 100
```

or

```
(gdb) l 1, 100
```

- Set a breakpoint at line 13:

```
(gdb) break 13
```

or

```
(gdb) b 13
```



GDB Usage Example

From within the GDB environment, let's issue the following commands (cont.):

- Run the program:

```
(gdb) run
```

or

```
(gdb) r
```

- Lets inspect the value of variable **x**:

```
(gdb) print x
```

or

```
(gdb) p x
```

- Execute line 14 (just one step):

```
(gdb) step
```

or

```
(gdb) s
```

- Lets inspect the value of variable **x** again:

```
(gdb) p x
```



GDB Usage Example

From within the GDB environment, let's issue the following commands (cont.):

- Which is the next line to be executed?

```
(gdb) frame
```

or

```
(gdb) f
```

- One more step:

```
(gdb) s
```

- Step over the call to **Average()**:

```
(gdb) next
```

or

```
(gdb) n
```

Differently from
step, next
do not step
into functions!

- Inspect the value of variable **avg**:

```
(gdb) p avg
```

It seems that the problem is within
the **Average()** function!



GDB Usage Example

From within the GDB environment, let's issue the following commands (cont.):

- Let's restart the program:

```
(gdb) r
```

- We've got stuck at line 13 again! First, let's print breakpoint information for the program:

```
(gdb) info breakpoint
```

or

```
(gdb) info b
```

- Now, delete breakpoint 1:

```
(gdb) delete 1
```

or

```
(gdb) d 1
```

- Step until we reach line 15.



GDB Usage Example

From within the GDB environment, let's issue the following commands (cont.):

- Now, let's step into the function **Average()**:

```
(gdb) s
```

→ step steps into functions!

- Let's check the value of the local variable **avg**:

```
(gdb) p avg
```

→ **avg was not properly initialized!**

- Initialize **avg** with 0!
- Now, recompile the program (do not close GDB!).
- Rerun the program, without breakpoints, and check the answer. → **It seems that we still have a problem!**



GDB Usage Example

From within the GDB environment, let's issue the following commands (cont.):

- Set a breakpoint at the function **Average()** :

```
(gdb) b Average
```

- Rerun the program (it will stop within **Average()**).
- Set a watchpoint for when the loop finishes (**i == n**):

```
(gdb) watch i == n
```

- Continue until next breakpoint / watchpoint.

```
(gdb) continue
```

or

```
(gdb) c
```



GDB Usage Example

From within the GDB environment, let's issue the following commands (cont.):

- Check the value of variables **avg** and **n**.
- Step until we leave **Average()**.
- Check the value that was returned by the function.



Damn! The value is incorrect! So what???

- The value returned by **Average()** is the result of a integer division! We have to cast one of the operators (e.g. **(float) avg/n**) to force a float division!
- Apply cast, recompile, delete all breakpoints and rerun!



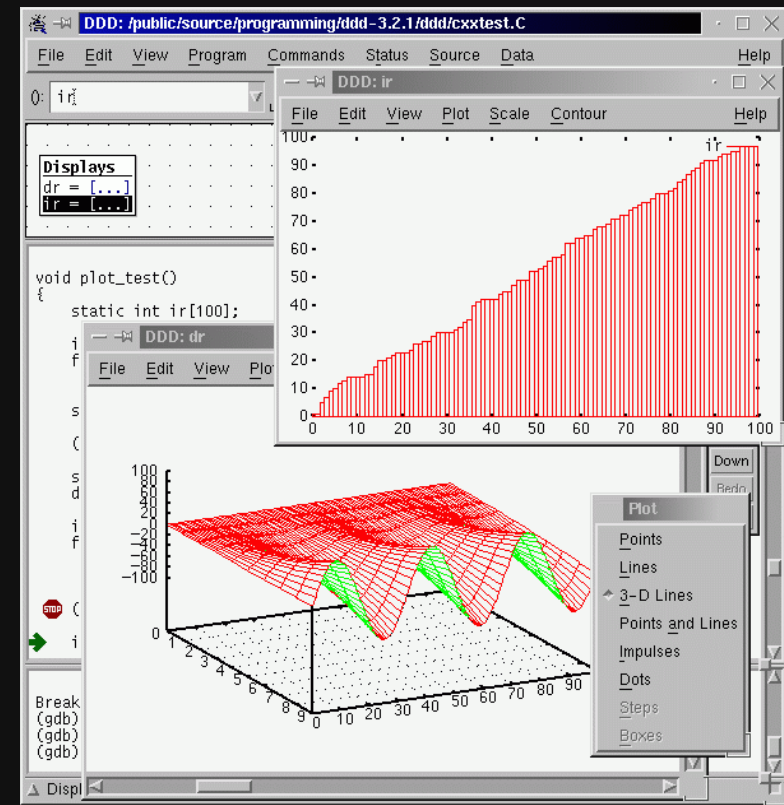
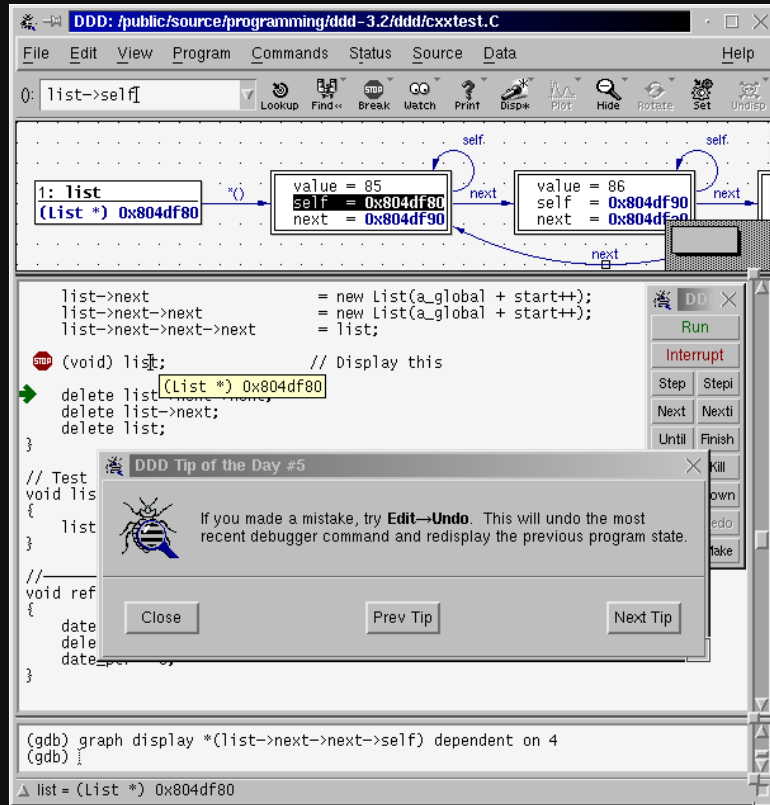
GDB Summary

- GDB is a quite powerful debugger tool.
- However, GDB does not allow one to easily follow the source code during a debug session.
- Some **GDB front ends** were developed, most notably:
 - CTRL + x + a : splits the GDB screen in command and source code windows (buggy!).
 - cgdb: curses based GDB front end.
 - Eclipse: an IDE that may use GDB as its debugging tool.
 - **DDD** : the Data Display Debugger.



The Data Display Debugger (DDD)

- It is a graphical interface to GDB.



- More on:
 - <https://www.gnu.org/software/ddd>



Back to C Data Types... Pointers!

A **pointer** variable is a **memory location** into which **data** (i.e. a **memory address**) can be **stored**.



Pointer Variable Example

Compile and run the following code from within GDB:

example_23.c

```
int main( void ) {  
    int x;  
    int *px;  
    x = 25;  
    px = &x;  
  
    return 0;  
}
```

00000000000000000000000000011001₂

byte

x

Addr.	Value
addr1 - 1	...
addr1 + 0	...
addr1 + 1	...
addr1 + 2	...
addr1 + 3	...

addr2 + 0	...
addr2 + 1	...
addr2 + 2	...
addr2 + 3	...
addr2 + 4	...
addr2 + 5	...
addr2 + 6	...
addr2 + 7	...

(gdb) b 4 —> Set a breakpoint at line 4.
(gdb) r —> Run (stops at line 4).
(gdb) p &x —> Print the address of x.
(gdb) p x —> Print the integer value stored at x.
(gdb) s

DIY!



Pointer Variable Example

Compile and run the following code from within GDB:

example_23.c

```
int main( void ) {
    int x;
    int *px;
    x = 25;
    px = &x;

    return 0;
}
```

00000000000000000000000000011001₂

00000000000000000111111111111111
111111111111111111101110100000100₂

byte

x

px

Addr.	Value
addr1 - 1	00011001
addr1 + 0	00000000
addr1 + 1	00000000
addr1 + 2	00000000
addr1 + 3	00000000

addr2 + 0	...
addr2 + 1	...
addr2 + 2	...
addr2 + 3	...
addr2 + 4	...
addr2 + 5	...
addr2 + 6	...
addr2 + 7	...

(gdb) b 4 —> Set a breakpoint at line 4.
 (gdb) r —> Run (stops at line 4).
 (gdb) p &x —> Print the address of x.
 (gdb) p x —> Print the integer value stored at x.
 (gdb) s —> Execute line 4 (x = 25).
 (gdb) p x —> Print the integer value stored at x.
 (gdb) p px —> Print the address stored at px.
 (gdb) s



Pointer Variable Example

Compile and run the following code from within GDB:

example_23.c

```
int main( void ) {
    int x;
    int *px;
    x = 25;
    px = &x;

    return 0;
}
```

00000000000000000000000000011001₂

00000000000000000111111111111111
111111111111111111101110100000100₂

byte

x

px

Addr.	Value
addr1 - 1	00011001
addr1 + 0	00000000
addr1 + 1	00000000
addr1 + 2	00000000
addr1 + 3	00000000

addr2 + 0	00000100
addr2 + 1	11011101
addr2 + 2	11111111
addr2 + 3	11111111
addr2 + 4	11111111
addr2 + 5	01111111
addr2 + 6	00000000
addr2 + 7	00000000

```
(gdb) b 4 —> Set a breakpoint at line 4.
(gdb) r —> Run (stops at line 4).
(gdb) p &x —> Print the address of x.
(gdb) p x —> Print the integer value stored at x.
(gdb) s —> Execute line 4 (x = 25).
(gdb) p x —> Print the integer value stored at x.
(gdb) p px —> Print the address stored at px.
(gdb) s —> Execute line 5 (px = &x).
(gdb) p px —> Print the address stored at px.
```



Examining Data with GDB: `print`

- `print` is the most **common** way to **examine data**, and is based on **expression evaluation**.
- `print` is able to **format the output!**
- Back to the previous example: set a breakpoint at line 7 (`return 0`) and run.

`example_23.c`

```
int main( void ) {  
    int x;  
    int *px;  
    x = 25;  
    px = &x;  
  
    return 0;  
}
```

Experiment `print` with these arguments:

<code>(gdb) p x</code>	→	Print the integer value stored at <code>x</code> .
<code>(gdb) p /x x</code>	→	Print the value at <code>x</code> in hexa format.
<code>(gdb) p /t x</code>	→	Print the value at <code>x</code> in binary format.



Examining Data with GDB: **x**

- The **x** command allows for **low-level** data examination.
- It prints the **contents** of **memory** positions in a specified **format**.



Examining Data with GDB: x

Back to the previous example:
set a breakpoint at line 7
(`return 0`) and run.

`example_23.c`

```
int main( void ) {  
    int x;  
    int *px;  
    x = 25;  
    px = &x;  
  
    return 0;  
}
```

byte →

Addr.	Value
<code>addr1 - 1</code>	00011001
<code>addr1 + 0</code>	00000000
<code>addr1 + 1</code>	00000000
<code>addr1 + 2</code>	00000000
<code>addr1 + 3</code>	00000000

`px`

<code>addr2 + 0</code>	00000100
<code>addr2 + 1</code>	11011101
<code>addr2 + 2</code>	11111111
<code>addr2 + 3</code>	11111111
<code>addr2 + 4</code>	11111111
<code>addr2 + 5</code>	01111111
<code>addr2 + 6</code>	00000000
<code>addr2 + 7</code>	00000000

Experiment `x` with these arguments:

(gdb) <code>x &x</code>	→	Print the value at address <code>&x</code> (last format).
(gdb) <code>x/t &x</code>	→	Print the value at address <code>&x</code> in binary.
(gdb) <code>x/d &x</code>	→	Print the value at address <code>&x</code> in decimal.
(gdb) <code>x/4tb &x</code>	→	Print 4 bytes in binary starting at addr. <code>&x</code> .
(gdb) <code>x/8tb &px</code>	→	Print 8 byt. in binary starting at addr. <code>&px</code> .



Scripting GDB

What if we would like to print the intermediary values of the summation below?

summation.c

```
#include <stdio.h>

int Sum( int begin, int end ) {
    int i;
    int acc = 0;
    for ( i = begin; i <= end; i++ )
        acc += i;

    return acc;
}

int main ( void ) {
    int a = 1;
    int b = 5;
    int sum = Sum( a, b );
    printf( "Sum: %i\n", sum );
    return 0;
}
```

We can automate it with **GDB scripting**, thus avoiding code modifications!



Scripting GDB

The following GDB script dumps on the screen all intermediary values generated during the summation computation:

`sumdebug.gdb`

Invoking GDB

```
~$ gdb --batch --command=sumdebug.gdb a.out
```

DIY!

```
set width 0
set height 0
set verbose off

b 8

commands 1
  silent
  printf "acc = %i\n", acc
  continue
end

b 10

commands 2
  silent
  printf "acc = %i\n", acc
  continue
end

run
```



Pointer Arithmetic

example_24.c

```
int x[4] = { 10, 20, 30, 40 };

int main( void ) {
    int *pint = x;

    *pint = 0;
    *(pint + 1) = 0;
    *(pint + 2) = 0;
    *(pint + 3) = 0;

    char *pbyte = ( char* ) x;

    *pbyte = 255;
    *(pbyte + 1) = 255;
    *(pbyte + 2) = 255;
    *(pbyte + 3) = 255;

    return 0;
}
```

pint →
pbyte → **x**
pbyte+1 →
pbyte+2 →
pbyte+3 →
pint+1 →

pint+2 →

byte →

Addr.	Value
addr + 0	00001010
addr + 1	00000000
addr + 2	00000000
addr + 3	00000000
addr + 4	00010100
addr + 5	00000000
addr + 6	00000000
addr + 7	00000000
addr + 8	00011110
addr + 9	00000000
addr + 10	00000000
addr + 11	00000000

■ ■ ■



Pointer Arithmetic

example_24.c

```
int x[4] = { 10, 20, 30, 40 };

int main( void ) {
    int *pint = x;

    *pint = 0;
    *(pint + 1) = 0;
    *(pint + 2) = 0;
    *(pint + 3) = 0;

    char *pbyte = ( char* ) x;

    *pbyte = 255;
    *(pbyte + 1) = 255;
    *(pbyte + 2) = 255;
    *(pbyte + 3) = 255;

    return 0;
}
```

Dereferencing
can be
equivalently
rewritten with []!

example_25.c

```
int x[4] = { 10, 20, 30, 40 };

int main( void ) {
    int *pint = x;

    pint[0] = 0;
    pint[1] = 0;
    pint[2] = 0;
    pint[3] = 0;

    char *pbyte = ( char* ) x;

    pbyte[0] = 255;
    pbyte[1] = 255;
    pbyte[2] = 255;
    pbyte[3] = 255;

    return 0;
}
```

Sounds
familiar?



Pointers vs. Arrays

- Are arrays **pointers**?
- If there is a difference, **could you point it out?**

Before discussing **how arrays** actually **work** in C, we will first take a look at **assembly!**



References

Learning C with GDB. Alan O' Donnell.

- <https://www.recurse.com/blog/5-learning-c-with-gdb>

